

Data Driven Improvements to RRT

15-780 Grad AI

Filipe Militão, Karl Naden, Bernardo Toninho

May 5, 2010

1 Introduction

The RRT algorithm represents an extreme in the design space of planning algorithms. It forsakes optimality and explores the space through randomness and a small bias towards goals. This is in marked contrast to other planning strategies such as visibility graphs or voronoi diagrams which carefully analyze and partition the space to find paths. This difference makes RRT very efficient even in higher dimensional spaces because its computational overhead is small in comparison to other algorithms which attempt to precisely characterize the world. However, this also means that RRT does not take advantage of information that it learns about the world while exploring.

The central idea of our project is to introduce learning of the environment into the RRT algorithm for online use in the search and also for re-planning purposes. We focus our efforts on the extension length for two reasons. First, the RRT algorithm gets information about good and bad extension lengths from given points when it tests to see if a given extension fails. Second, the optimal extension length varies based on the obstacles surrounding a given point. The goal of this project is to extend the RRT planning and the ERRT re-planning algorithm to store and use information learned about good extension lengths while searching the space. Our hope is that we can realize significant improvements in how fast and often RRT-based planning finds the goal, without sacrificing its key principles.

Our main goal is to explore a set of variations in the standard RRT (and ERRT) algorithms to try

and offer a better algorithm for these heterogeneous worlds both in planning and re-planning scenarios.

Some other papers describing extensions to the RRT algorithm include [Lav98], [LL04], [JYLS05], and [YL09].

2 RRT Algorithms

A Rapidly-exploring Random Tree (RRT) algorithm is a random algorithm designed for efficient search in continuous spaces. The algorithm finds a path from a start point S to a goal G in a given search space by generating a random tree in the space. The tree will have S as its root and is computed by randomly choosing points in the world, finding the tree node N closest to the random point and extending the tree by creating a new node N' connected to N , in the direction of the random point (as long as the extension does not intersect any obstacles in the world).

An immediate optimization to the algorithm in terms of path planning, called RRT-GoalBias, consists of adding a probability parameter p , which causes the algorithm to choose to extend the tree towards the goal with probability p , and extend towards a random point with probability $1 - p$.

The main advantages of using a RRT algorithm are that it does not require a discretization of the space and the high efficiency in terms of time and space usage. Furthermore, it turns out that such a biased random exploration of the space does indeed prove to be successful in finding paths, even in the presence of a large number of obstacles. Such pos-

itive results have resulted in the usage of RRT for robot motion planning and re-planning (in the latter setting, information about previous paths is also used to bias the search).

2.1 Context Sensitivity

Consider now a realistic scenario where the space in which we desire to plan has a high number of obstacles, possibly non-uniformly distributed throughout the world. Such a world will have regions which will be fairly open, in which should be easy to find a path, and regions which will be more cluttered and therefore harder to navigate through. The way RRT is used to find paths in such worlds, simply relies on lowering the probability p and the extension length. The first aims to avoid being stuck in the obstructed regions of the world, while the latter aims to allow for extensions to succeed in these regions. Unfortunately, lowering these parameters naturally decreases overall performance, since the algorithm extends less often towards the goal, and requires more extensions (and therefore more computation) to actually reach it, even in regions of the space where such a conservative approach may not be required.

Our observation is that RRT does not take advantage of the information it collects about the world, and therefore always employs a rigid exploration tactic that needs to be globally conservative to succeed in the scenario above. However, it should be possible to use the information collected throughout the exploration to adapt the tree generation procedure in such a way that the tree will cover the less obstructed regions of space faster (in terms of computation time and number of nodes), all the while maintaining the ability to navigate through the more obstructed areas. We address this issue by allowing the extension lengths between nodes to change over the course of the tree generation in such a way that the extension length becomes higher in open areas and lower in cluttered ones. Our basic framework is as follows: we associate with each node in the tree an extension factor, which is used to determine the extension length that is to be applied to all extensions from the node. Whenever an extension from a node fails, we decrease its extension factor. When

an extension succeeds, we increase the node's extension factor, generate the new node which will inherit the extension factor of its parent.

The idea is that we should always try to eagerly maximize the extension length, in order to cover space as fast as possible, therefore as long as extensions succeed (and therefore the algorithm does not run into obstacles) it keeps increasing the extension length accordingly. When the algorithm cannot extend from a node, it means that an obstacle exists in the proximity of that node and therefore the extension length should decrease accordingly, to potentially increase the success of navigating through.

2.1.1 VLRRT

Variable Length RRT (VLRRT) is an algorithm that follows immediately in our framework. With each node of the tree, we associate a value ϵ which is used as a multiplicative factor to the base extension length. For the root of the tree, the value of ϵ starts off as 1. As in standard RRT-GoalBias, the algorithm is parametrized with a probability of extending towards the goal. The main difference lies in the treatment of the actual extensions. Whenever an extension from a given node succeeds, we increase its ϵ value and generate the new node with the new ϵ value of its parent, in order to propagate the extension information throughout the tree. Whenever an extension fails, we decrease the ϵ value of the node.

A natural question arises in this scenario, which is how to vary the ϵ value in the presence of successful and unsuccessful extensions. Therefore, we opted to parametrize the algorithm with three possible schemes with which to increase and/or decrease ϵ : a *linear* scheme, in which ϵ increases/decreases a fixed amount; a *multiplicative* scheme, in which ϵ is multiplied or divided by a fixed amount, and a *restart* scheme, in which ϵ is restarted to its initial value of 1.

2.1.2 DVLRRRT

An immediate observation regarding VLRRT is that a successful or unsuccessful extension provides us information of the presence or absence of obstacles

in a particular direction, but not in all directions (which is the generalization that VLRRT employs). While such an over-generalization may indeed provide good results (which indeed turns out to be the case), it is clear that there will be scenarios where ϵ will increase/decrease when perhaps it didn't need to. To take the extension idea one step further, we developed Directional VLRRT (DVLRRRT).

The idea behind DVLRRRT is to account directionality into VLRRT's ϵ value modifications, in order to overcome the potential downsides of the generalization employed in VLRRT. In each tree node, we store a *directional ϵ map*. This structure maps a particular direction with an epsilon value, which enables us to take into account directions when updating ϵ and computing extension lengths.

Whenever we attempt to extend from a node in a given direction d , we check if a mapping with that particular direction exists in the directional map. If so, we use the epsilon value stored in the map to compute the extension length and update it on failure/success as in VLRRT. If the mapping does not yet exist, we compute the ϵ value for the given direction by taking into account the directional information in the map: we look up directions similar to d up to a certain angular distance, clockwise and counter-clockwise, and weight the information by proximity to d . We update the directional map with the ϵ value for d on extension success/failure accordingly.

With this new strategy, we inform our extension length variation even further, by taking into account extension success and failure in directions that are similar to ones we have observed before in a node.

3 Evaluating the algorithms

To evaluate the performance of our algorithms, we implemented a simulation suite (with visualization) that allows us to compare DVLRRRT, VLRRT and RRT-GoalBias. The suite is also prepared for replanning simulation, by simulating random changes in worlds in successive iterations.

Our testing suite executes the specified algorithms (with a given time limit for each execution) in a number of previously generated worlds and records

relevant statistics such as time spent to find the goal, path length, world coverage, etc.

The metrics we used to compare DVLRRRT, VLRRT and RRT-GoalBias for this report are success rate (in terms of finding the goal) and running time. For DVLRRRT and VLRRT, we also analyzed how the several ϵ modification schemes affect overall performance.

4 Planning Results

We executed each algorithm 10000 times, in several representative worlds, with a time limit of 50 milliseconds : a maze-like world *maze* (similar to the one in [BV02]), a largely unobstructed world with a particularly obstructed region in the path to the goal *obstructed* and a more highly cluttered world *cluttered*.

We can observe that in all the test worlds, our algorithms perform generally better (or nor worse) than RRT-GoalBias, achieving higher average success rates and lower average execution times, noting that in worlds where RRT-GoalBias already performs very well, the gains achieved by our algorithms are naturally smaller.

However, and not unexpectedly, the performance increase of our algorithms does not occur for all ϵ modification schemes. Our intuition was that we wanted to "speed through" what we consider to be unobstructed regions, but decrease our "speed" as fast as possible when we find an obstructed region, so we can effectively navigate through it. As it turns out, our intuition is indeed correct, seeing as the modification schemes that perform the best are those that increase ϵ faster but decrease ϵ even faster, in both VLRRT and DVLRRRT. Particularly, the best results were obtained when the increase rate was the highest (multiplicative scheme) and the decrease rate was the highest as well (restarting scheme). For the sake of readability, we omit the results regarding other modification schemes, noting that, in general, choosing a "slow" decrease rate produces overall bad results.

Analyzing the actual values (Fig. 1), we can see that in the *obstructed* world, RRT-GoalBias already

Algorithm	World	Avg. Success	Avg. Time (ms.)
RRT-GoalBias	obstr.	95%	11.94
VLRRT (M2/R)	obstr.	96%	13.48
VLRRT (M4/R)	obstr.	95%	13.32
DVLRRT (M2/R1)	obstr.	96%	13.65
DVLRRT (M3/R1)	obstr.	96%	13.41
DVLRRT (M4/R1)	obstr.	96%	13.38
VLRRT (M3/R1)	obstr.	95%	13.61
RRT-GoalBias	maze	91%	17.26
VLRRT (M2/R)	maze	99%	10.59
VLRRT (M4/R)	maze	99%	10.60
DVLRRT (M2/R1)	maze	99%	10.50
DVLRRT (M3/R1)	maze	99%	10.69
DVLRRT (M4/R1)	maze	99%	10.74
VLRRT (M3/R1)	maze	99%	10.69
RRT-GoalBias	clutt.	33%	39.26
VLRRT (M2/R)	clutt.	60%	34.00
VLRRT (M4/R)	clutt.	59%	33.61
DVLRRT (M2/R1)	clutt.	60%	33.86
DVLRRT (M3/R1)	clutt.	60%	33.73
DVLRRT (M4/R1)	clutt.	60%	33.89
VLRRT (M3/R1)	clutt.	60%	33.90

Figure 1: Success Rates and Exec. Times

performs very well, therefore our gains are less considerable. In fact, we gain a marginal overhead in terms of time due to extra computation. In the *maze* world, RRT-GoalBias performs a bit worse, since the world is generally more complex. In this world, our algorithms start to shine. Not only do they find they goal almost 100% of the runs, they actually find the goal faster. This shows that our algorithms are indeed able to navigate such worlds faster and more successfully, given that they can adapt themselves to the world in question. Finally, in the *cluttered* world, RRT-GoalBias does not fair well at all, given the high degree of heterogeneity in the world. Our algorithms however (not unexpectedly), manage to not only find paths to the goal more frequently (approx. twice as better), but also do so faster. Overall, we can conclude that our algorithms obtain an increase of success rate (in terms of finding the goal) of approximately 16% and a decrease in running time of approximately 15%, depending on the type of world we consider.

Algorithm	Avg. Succ. Rate	Avg. Time
VLRRT (M2/R)	17%	-15%
VLRRT (M4/R)	16%	-16%
DVLRRT (M2/R1)	17%	-15%
DVLRRT (M3/R1)	16%	-16%
DVLRRT (M4/R1)	17%	-15%
VLRRT (M3/R1)	16%	-15%

Figure 2: Overall Gains versus RRT

5 Re-planning

We now are in a situation where the world has changed slightly w.r.t the previous planning iteration. Our goal is to attempt to extract as much useful information as possible from what was gathered by the previous search, in the hope of reaching the goal faster than by just re-starting the search from scratch for the new world configuration. However, we have to be very frugal with computation in order to make our abstraction worthwhile. If time were no object, we would simply test the world to get perfect information. Thus, the challenge is to extract the pieces of information that will assist in re-planning without spending much time analyzing or refining the data.

For this re-planning component we base our algorithms on ERRT [BV02] that extends RRT with randomly extracted waypoints from the previous best path to the goal. Thus, all our variants were extended to include waypoints in the same fashion as ERRT.

With this in mind, we defined the following re-planning strategies:

Estimate a good initial step Use the previous random tree to learn a potentially better initial extension length for each new node of the search. Each estimation is customized for a specific algorithm:

(VL) Simple average of ϵ values from the K nearest neighbors of the previous tree;

(VL) Weighted average (based on the distance to the new node) of ϵ values from the K nearest neighbors;

(DVL) Pull the directional information from the nearest neighbors (adjusting the direction appropriately).

Biased waypoint generation The ERRT algorithm uses random nodes from the previous best path to goal as waypoints. We attempted to improve this by biasing the waypoint selection to nodes that are in less explored areas (lower densities) of the previous best path to goal. The intuition was that these would potentially be more valuable since an area being less explored usually means that it has more obstacles and is therefore harder to navigate through. However, empirical results (even without changing the world) actually show some performance degradation or no significant improvement at all and thus we will not show these results in detail due to space constraints.

6 Re-planning Results

For re-planning, we used the runtime average of 1000 test runs each with 100 planning iterations where the start point is slightly advanced towards the goal (using the path that got nearest to it, if none reached it) each with the same 50 ms time slice as before, and causing small changes in the world (Brownian motion - small random shifts in the position of the obstacles). We tested the performance of all re-planning algorithms, and our initial planning algorithms (basic VL and DVLRRRT).

Our results (Figs. 6, 7, 8) show three groups: one for ERRT, another for the VL variants (both for planning and re-planning) and the final one for both DVL alternatives (also both for planning and re-planning). Thus, we can conclude that there were no significant gains from using the more complex learning techniques (that try to compute a better initial step from the previous tree) when compared with just using the traditional (planning) versions, given that the average runtimes are virtually identical. As before, both our variants do rather well when compared to ERRT on most worlds, although the overhead of DVL is considerable when compared to VL.

Overall, this shows that the online learning techniques employed for VLRRT are sufficiently advanced to cope with this kind of Brownian motion in the world and capable of extracting valuable information (by testing) much more efficiently than any of our re-planning improvements.

7 Conclusions

We have shown that we can significantly reduce planning time and increase robustness by employing online learning techniques to RRT, at the cost of a sometimes slightly increased path length. Directional information proved not to produce meaningful benefits (most likely due to the small number of extensions applied to each node), while a simple combination of multiplicative increases combined with a step reset on extension failure was the alternative with best results. However, reusing this data for re-planning proved to be more difficult: there were no significant performance gains from computing estimations of the new world based on the previous information when compared to just consulting the world (i.e. running the planning algorithms from scratch). This also shows the flexibility of our regular planning techniques that are able to adjust very quickly to new world configurations.

As a final note, all code is publicly available (under an open-source license) at:


<http://code.google.com/p/vlerrt/>

References


- [BV02] James Bruce and Manuela Veloso. Real-time randomized path planning for robot navigation, 2002.
- [CH67] T. Cover and P. Hart. Nearest neighbor pattern classification. *Information Theory, IEEE Transactions on*, 13(1):21–27, 1967.
- [JYLS05] Lonard Jaillet, Anna Yershova, Steven M. LaValle, and Thierry Simon. Adaptive tuning of the sampling domain for dynamic-domain rrts, 2005.

-

A World Images



B Results



MAY 5, 2010




Figure 7: Obstructed re-planning.




Figure 8: Cluttered re-planning.